Hadoop Overview for Managers

PART 1

Outline

PART 1

- Hadoop Overview
  - Traditional Computing Systems and Limitations
  - Why Big Data?
  - Why Hadoop?
  - Hadoop Basic Concepts
  - Where Hadoop fits in the Enterprise
- Hadoop Architecture
  - Building blocks
  - HDFS
  - Demo
- **Break**

PART 2

- YARN Architecture
  - Yarn Overview
  - MapReduce
  - Demo
- Tools and technology for Hadoop ecosystem
- Hadoop Real Life Use Cases
- Establishing a Big Data Center of Excellence
  - Justifying business value for your organization
  - Challenges on building a production solution
  - Recommended organizational structure
  - Best Practices: Steps to effectively deploy Hadoop
- Recap and Q&A

ERA OF BIG DATA

Era of Big-Data

Source: Microsoft TechEd North America 2014

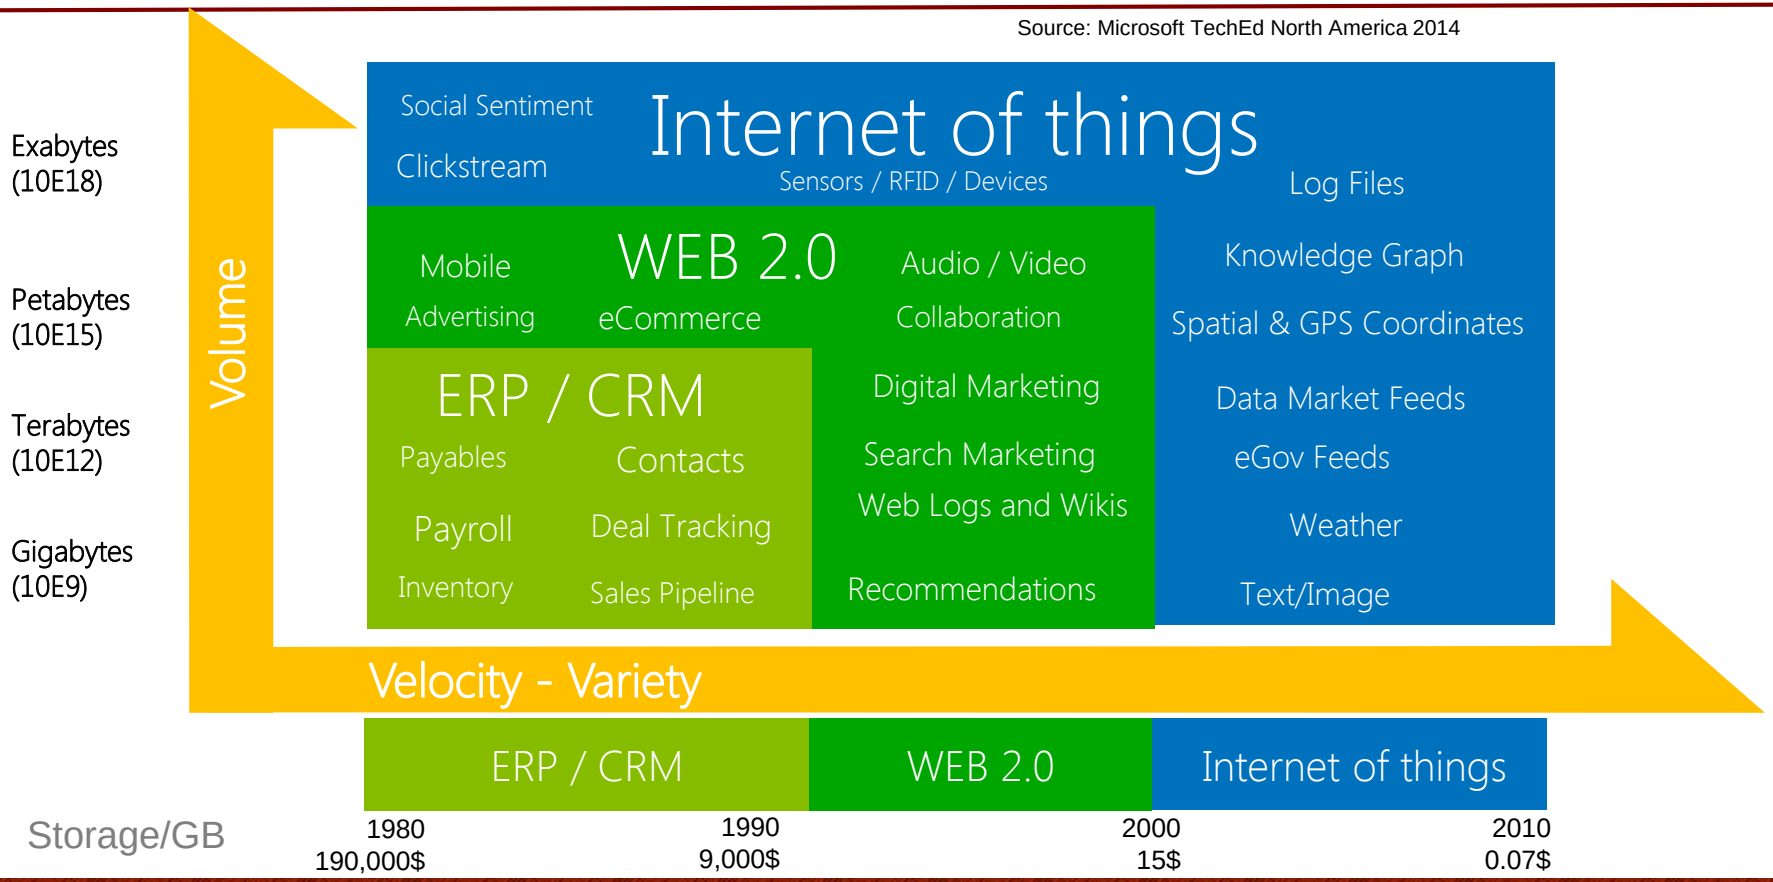

Who is generating so much data?

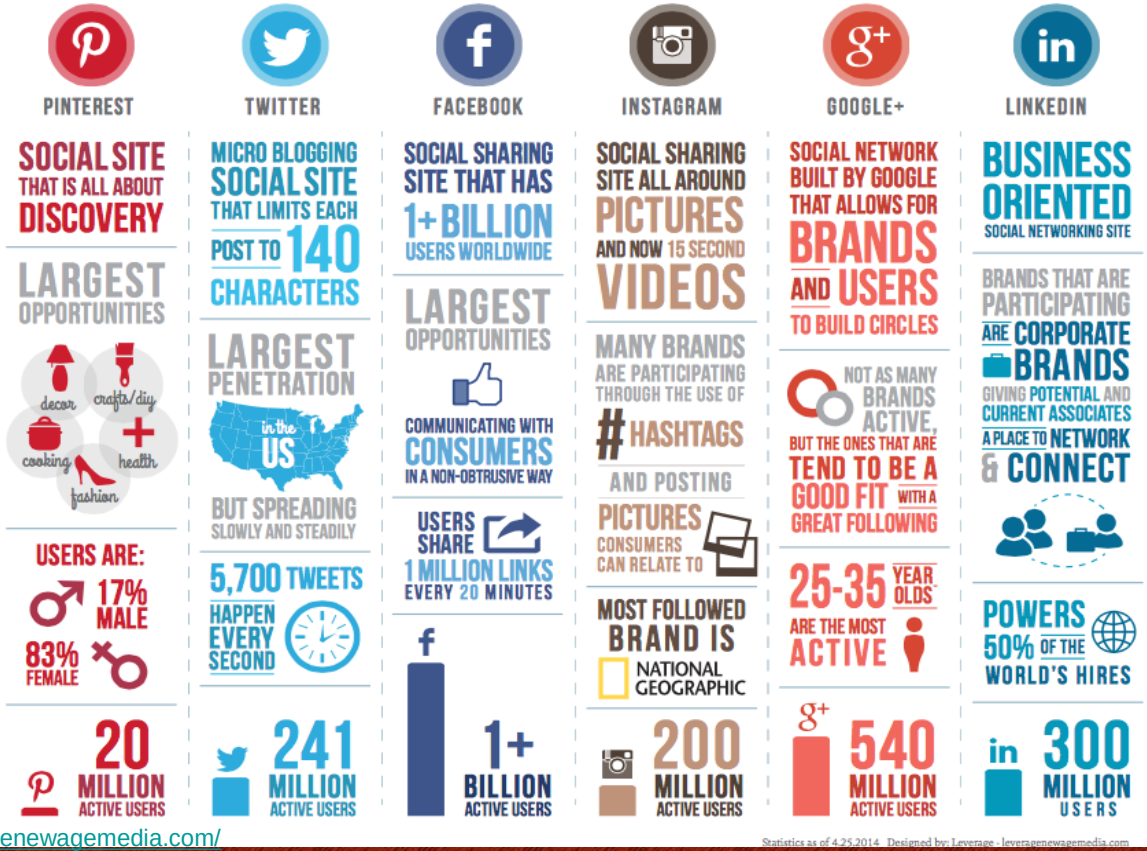

Source https://leveragenewagemedia.com/

Big Data in Healthcare - Asthmapolis

Collects data from patients (inhalers) and using analytics help them better manage their Asthma

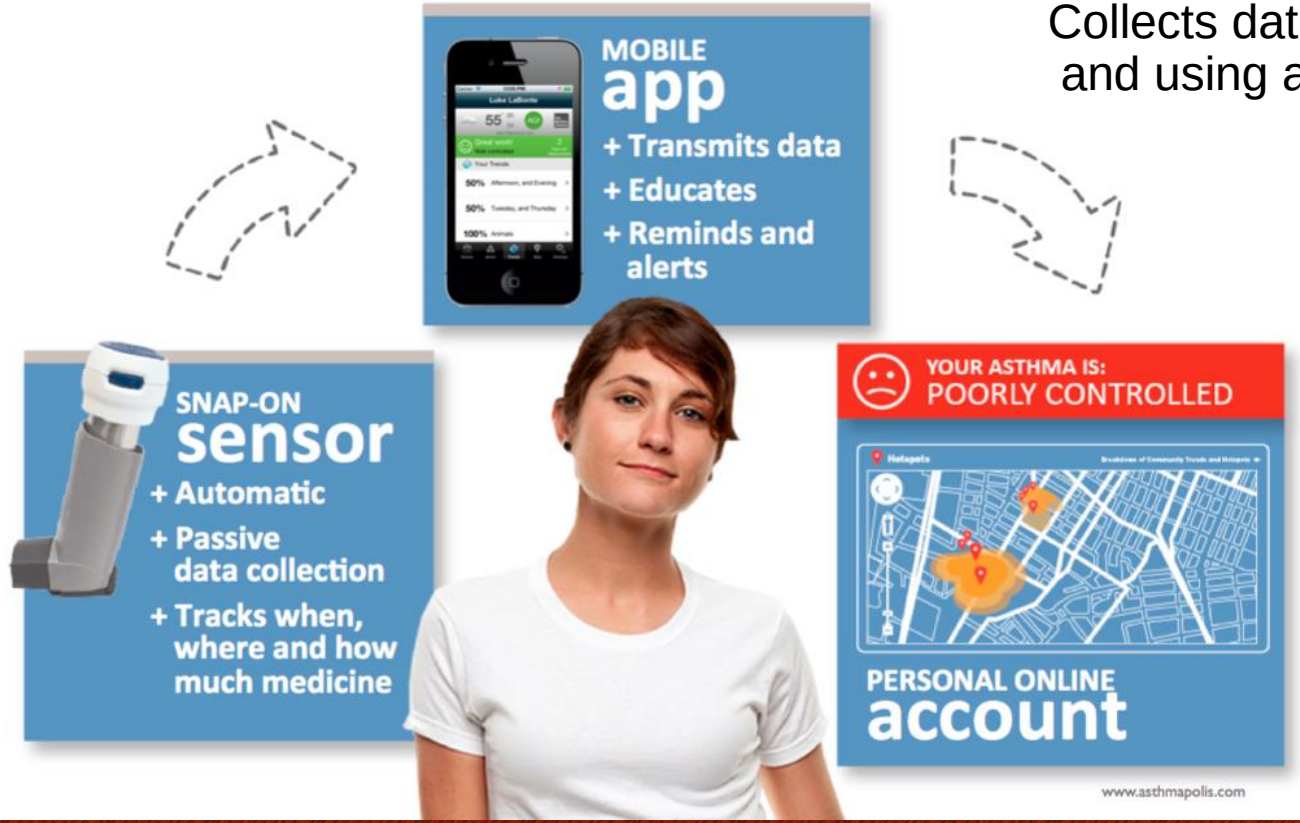

Big Data Solution – key characteristics

- **Must Scale with increasing volume**
  - Performance
  - Availability
  - Cost
- **Support variety of data**
  - Structured
  - Unstructured
  - Semi Structured

“Shared Everything” Architecture

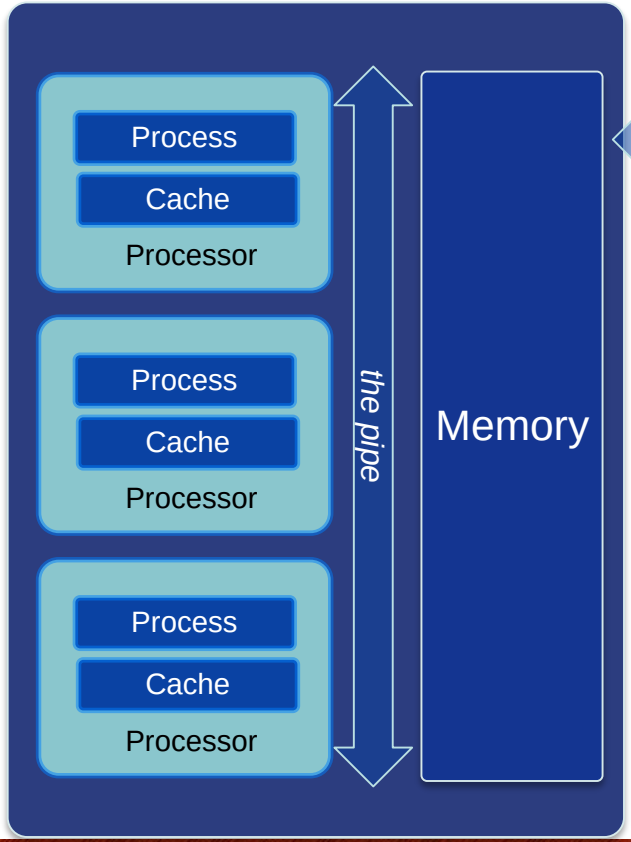

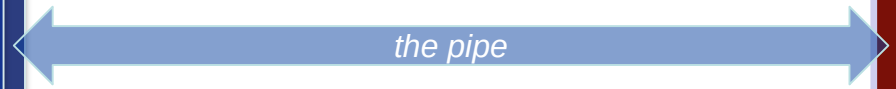

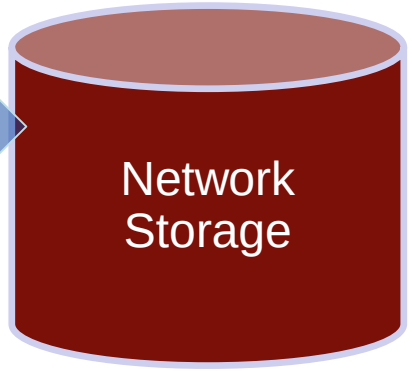

- Scale Up Architecture
- All Processors share same Memory Address Space
- Sharing same system bus may result in choking of the bandwidth for memory access
- Complex synchronization

(Is) Distributed Systems the solution

ZALONI

- Programming Model is complex
- Data exchange requires synchronization
- Failures are expensive and needs to be managed
- Does not scale for Large volumes of data – network interconnects in a datacenter are expensive!

New Approach to Distributed Systems

Must Scale with increasing volume
- Performance
- Availability
- Cost

Support variety of data
- Structured
- Unstructured
- Semi Structured

- Shared Nothing Architecture
- **Data Locality**
- No synchronization requirement among the nodes
- **Designed for failure** – Multiple copies of data
- Consistent – individual failures does not fail the job
- **Support "commodity"** hardware & heterogeneous

Shared Nothing Architecture

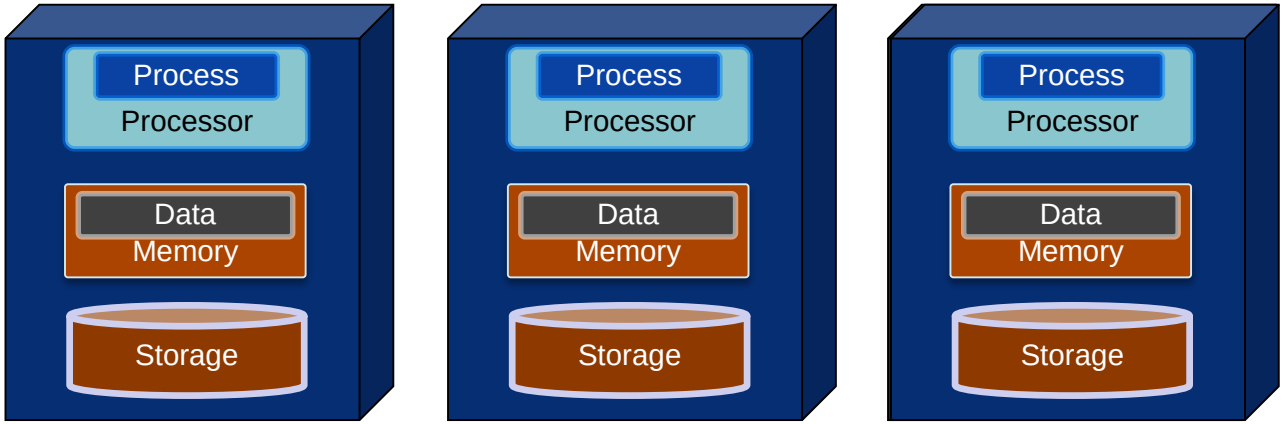

- Data is sharded (partitioned) amongst the nodes
- Computation is local to the nodes – no need to get the data from elsewhere
- No synchronization, Simple implementation

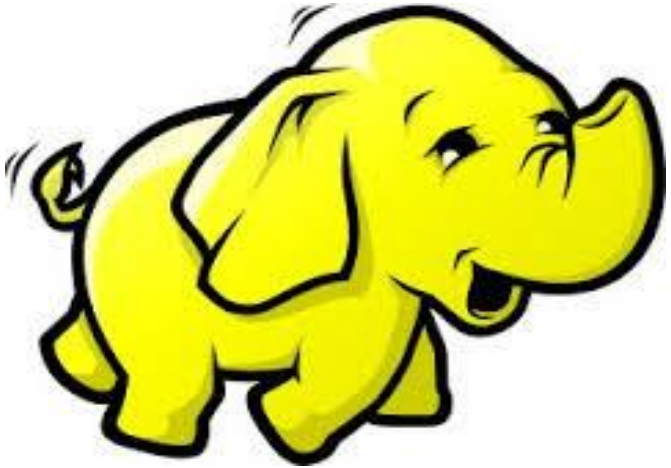

HADOOP OVERVIEW

Hadoop Core Components

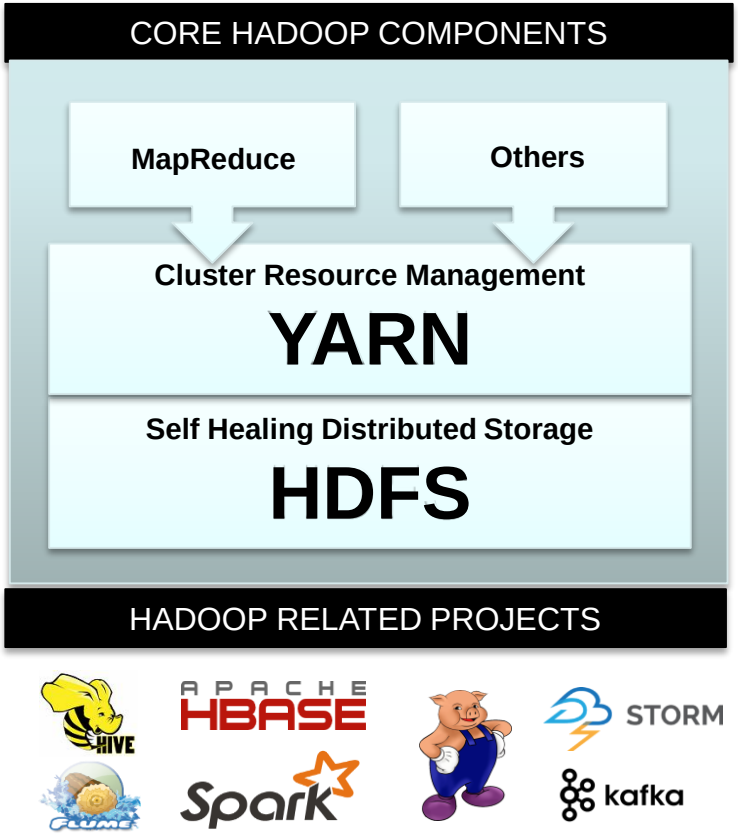

- **Hadoop Common:** A set of common libraries and utilities used by Hadoop modules.
- **Hadoop Distributed File System (HDFS)**: A scalable and fault tolerant distributed filesystem to data in any form.
- **Yet Another Resource Negotiator (YARN)**: From Hadoop 2.0, YARN is the cluster management layer to handle various workloads on the cluster.
- **MapReduce:** MapReduce is a framework that allows parallel processing of data in Hadoop.

Growing number of eco-system Projects

HDFS Overview

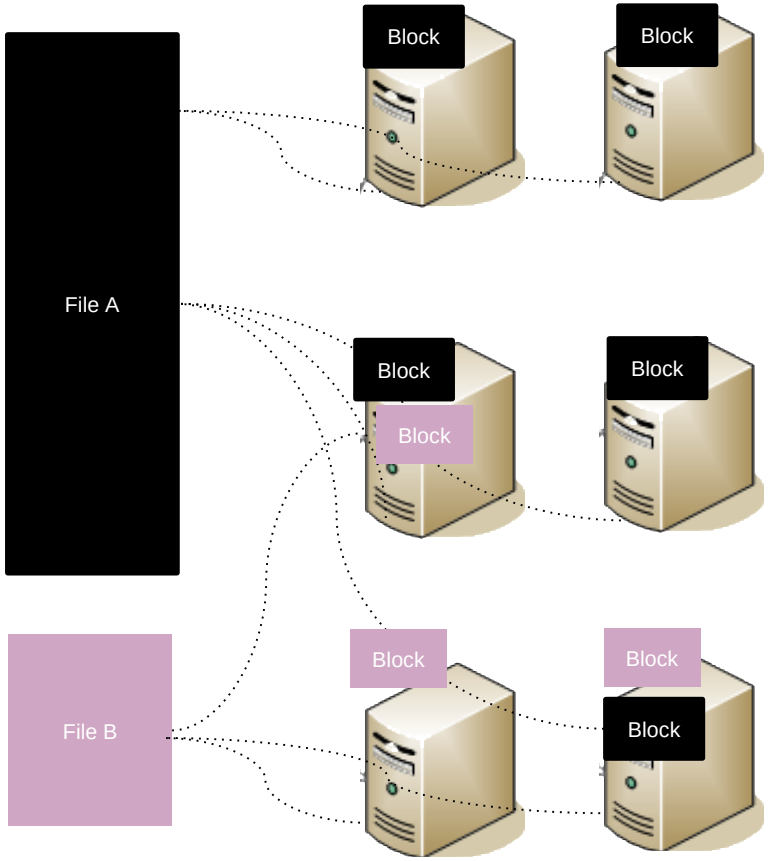

- **Distributed Storage**

- Designed to store very large files
- Sharded storage for high throughput
- Replicated storage for failure protection
- Self healing
- Add new disks or nodes to scale

YARN Overview

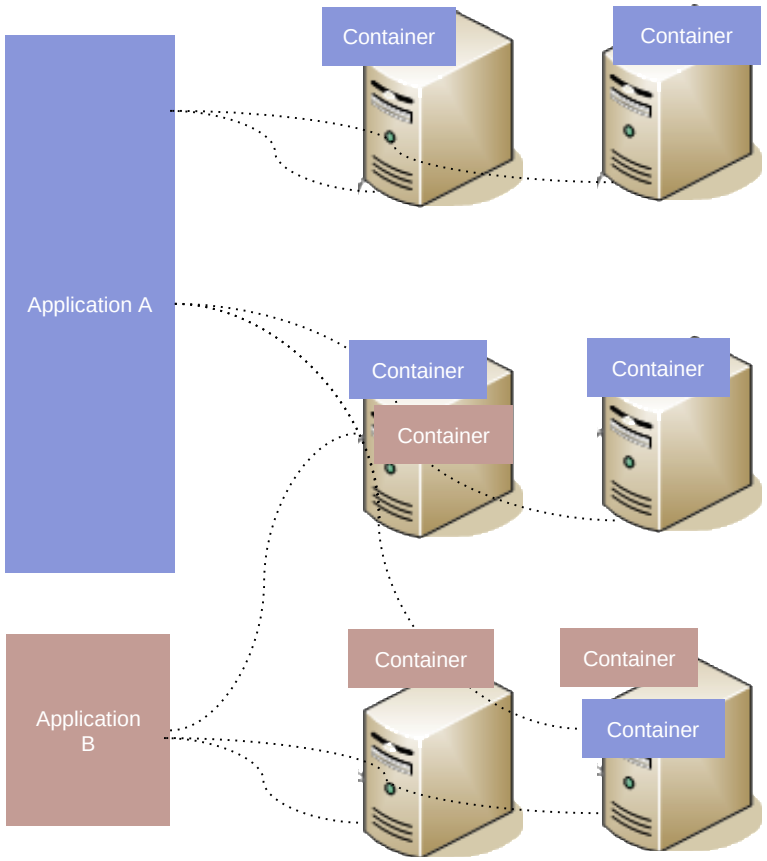

- **Distributed Compute**
- Yet Another Resource Negotiator
- Run Java code on arbitrary node(s) depending on availability
- “Cloud without the Virtualization”
- Applications can occupy resources as per need.
- Common model of distributing code and accessing Data
- Self healing

MapReduce Overview

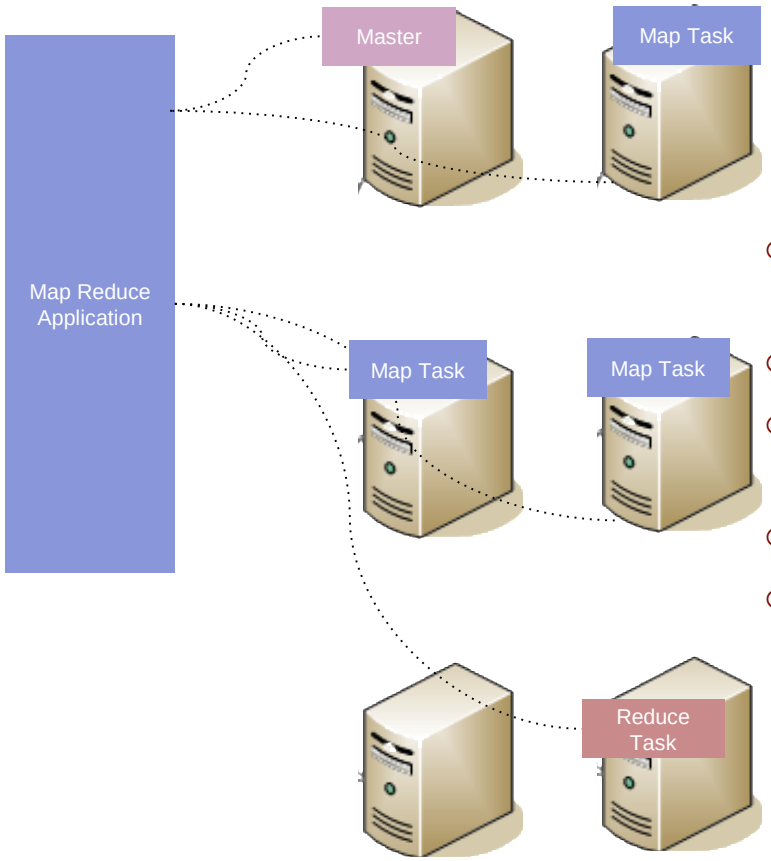

- Model for processing large amount of data in parallel
- Oriented towards batch processing
- Programming Model derived from functional programming
- Built on top of YARN
- I/O on HDFS and others

Timeline

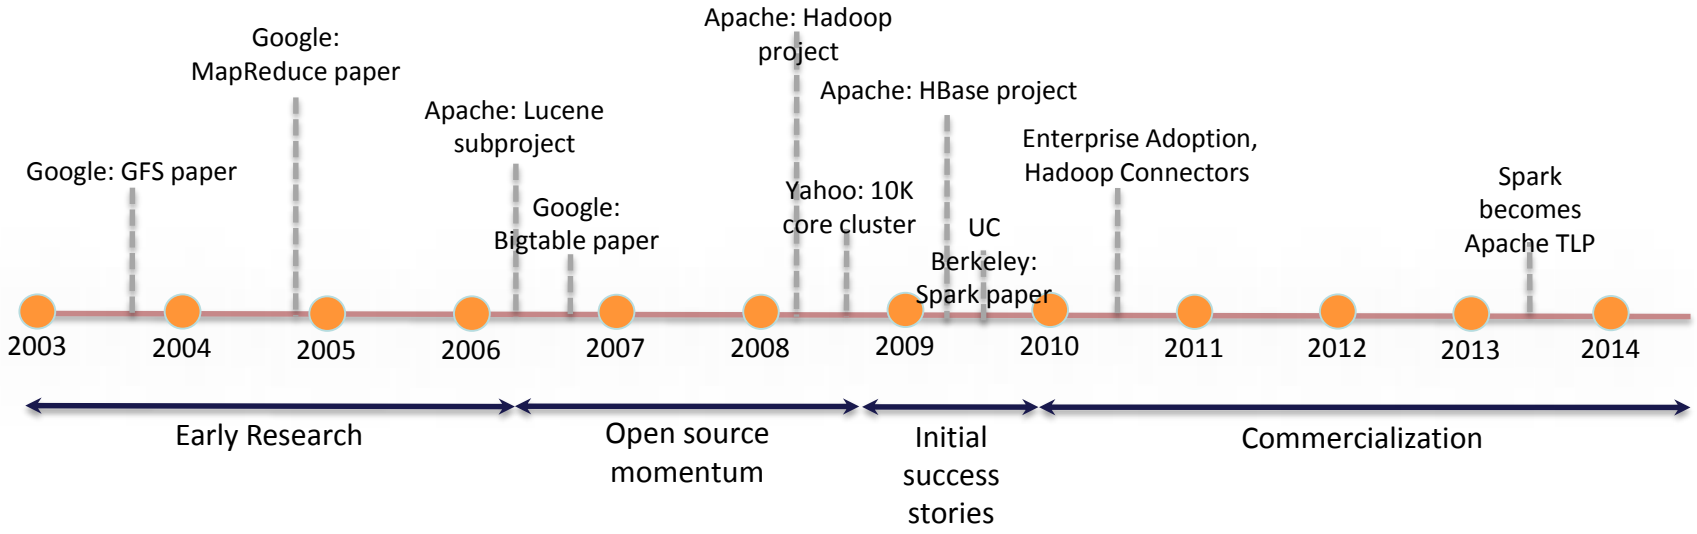

Why Hadoop?

- **Runs on commodity hardware (and the cloud)**
  - Low cost
  - Ease of maintenance
- Scales well
- **Strong Ecosystem**
- **Open Source**
  - Apache 2.0 License
  - Strong Community
- Runs on the JVM

When to use

Relational Databases:

Use when:

- Interactive OLAP Analytics (<1sec)
- Multistep ACID Transactions
- 100% SQL Compliance

Hadoop:

Use when:

- Structured or Not (Flexibility)
- Scalability of Storage/Compute
- Complex Data Processing

HADOOP DISTRIBUTION AND FRAMEWORKS

Cloudera CDH

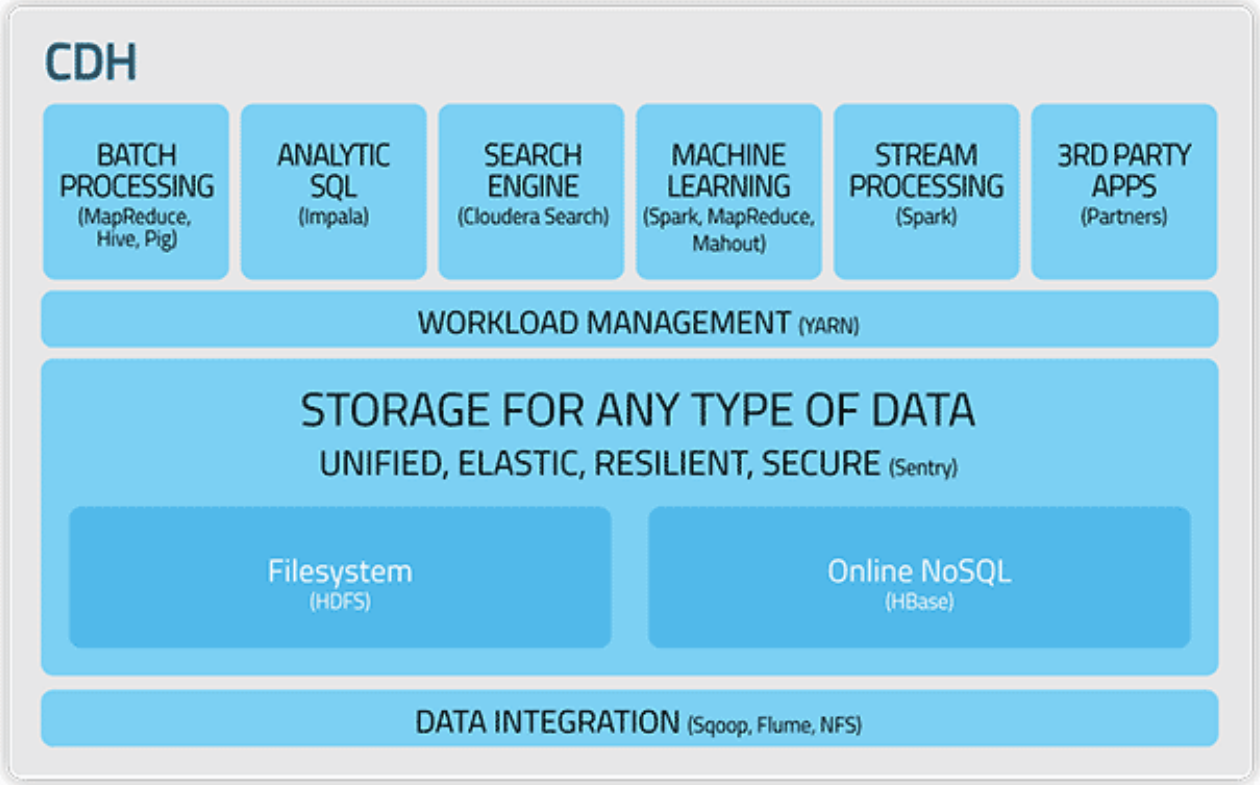

Hortonworks HDP

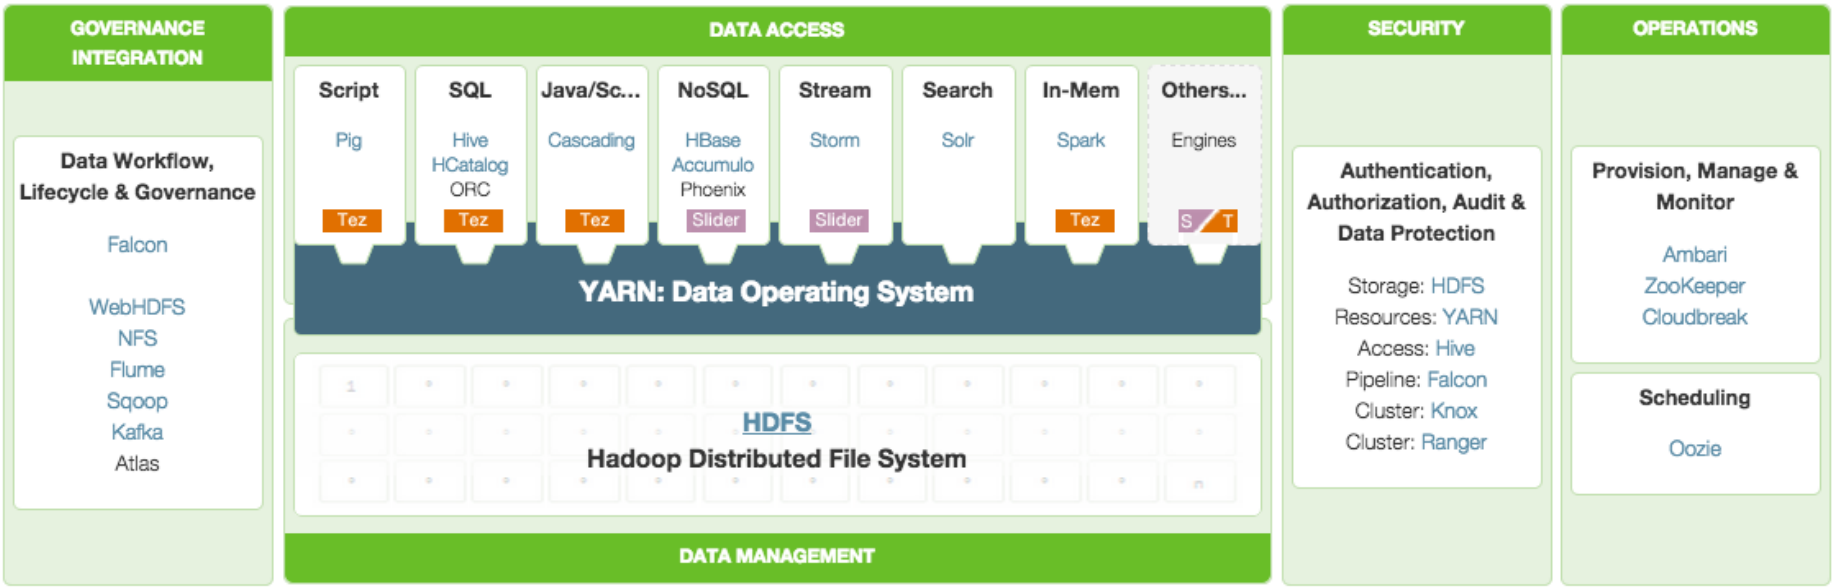

MapR M5

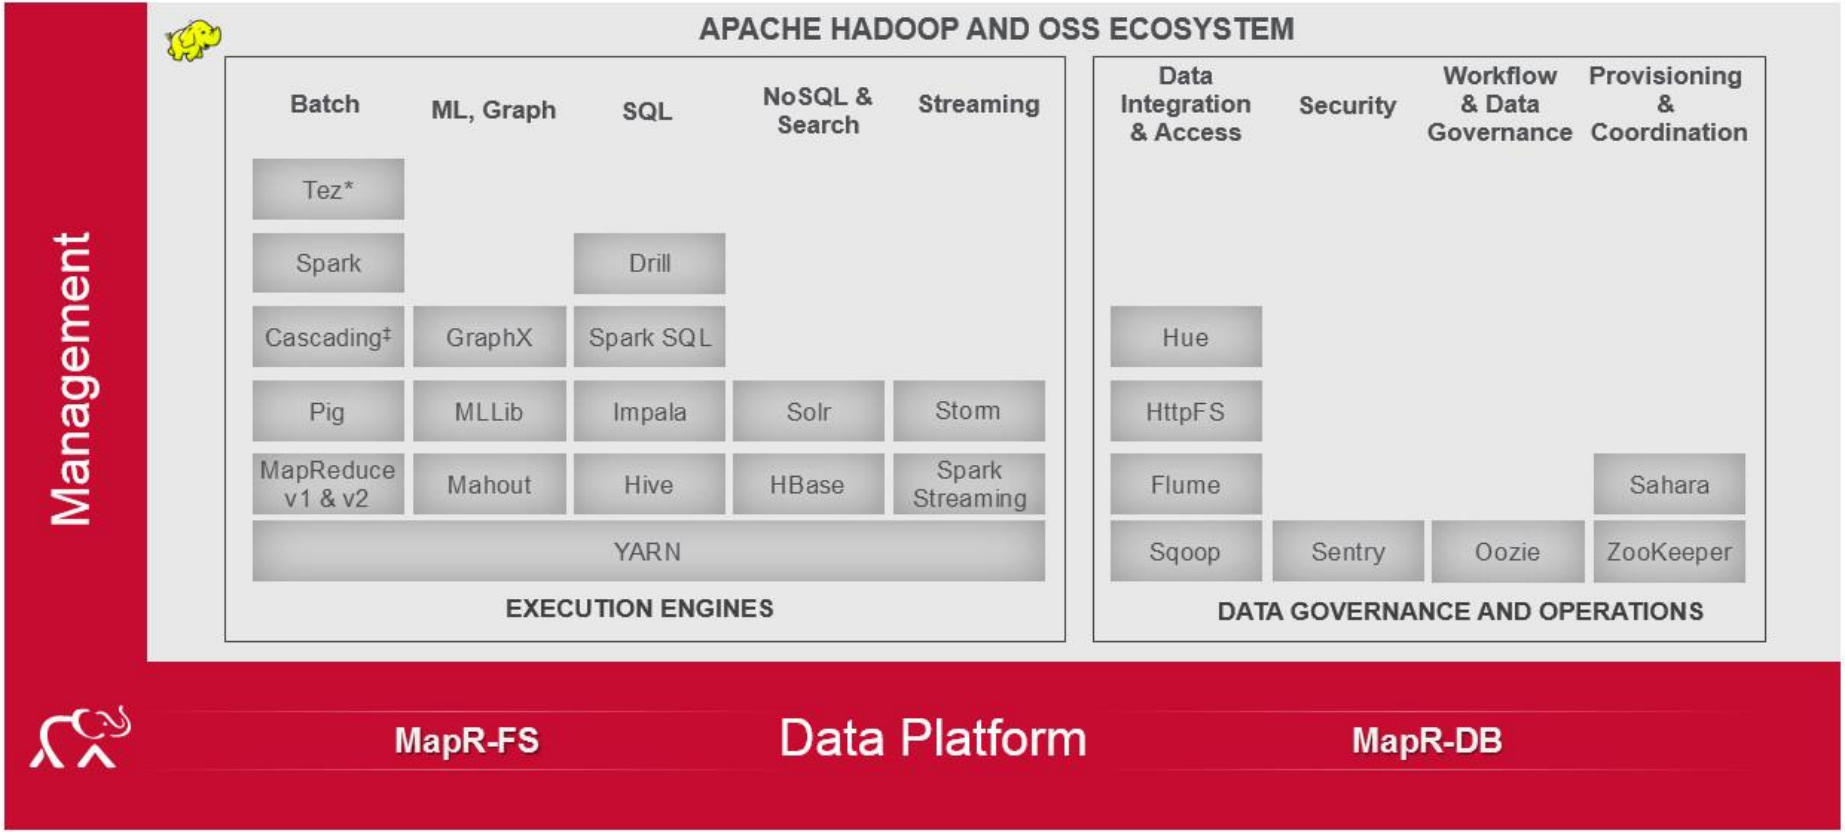

*Developer preview ‡Certified on MapR

Pivotal HD

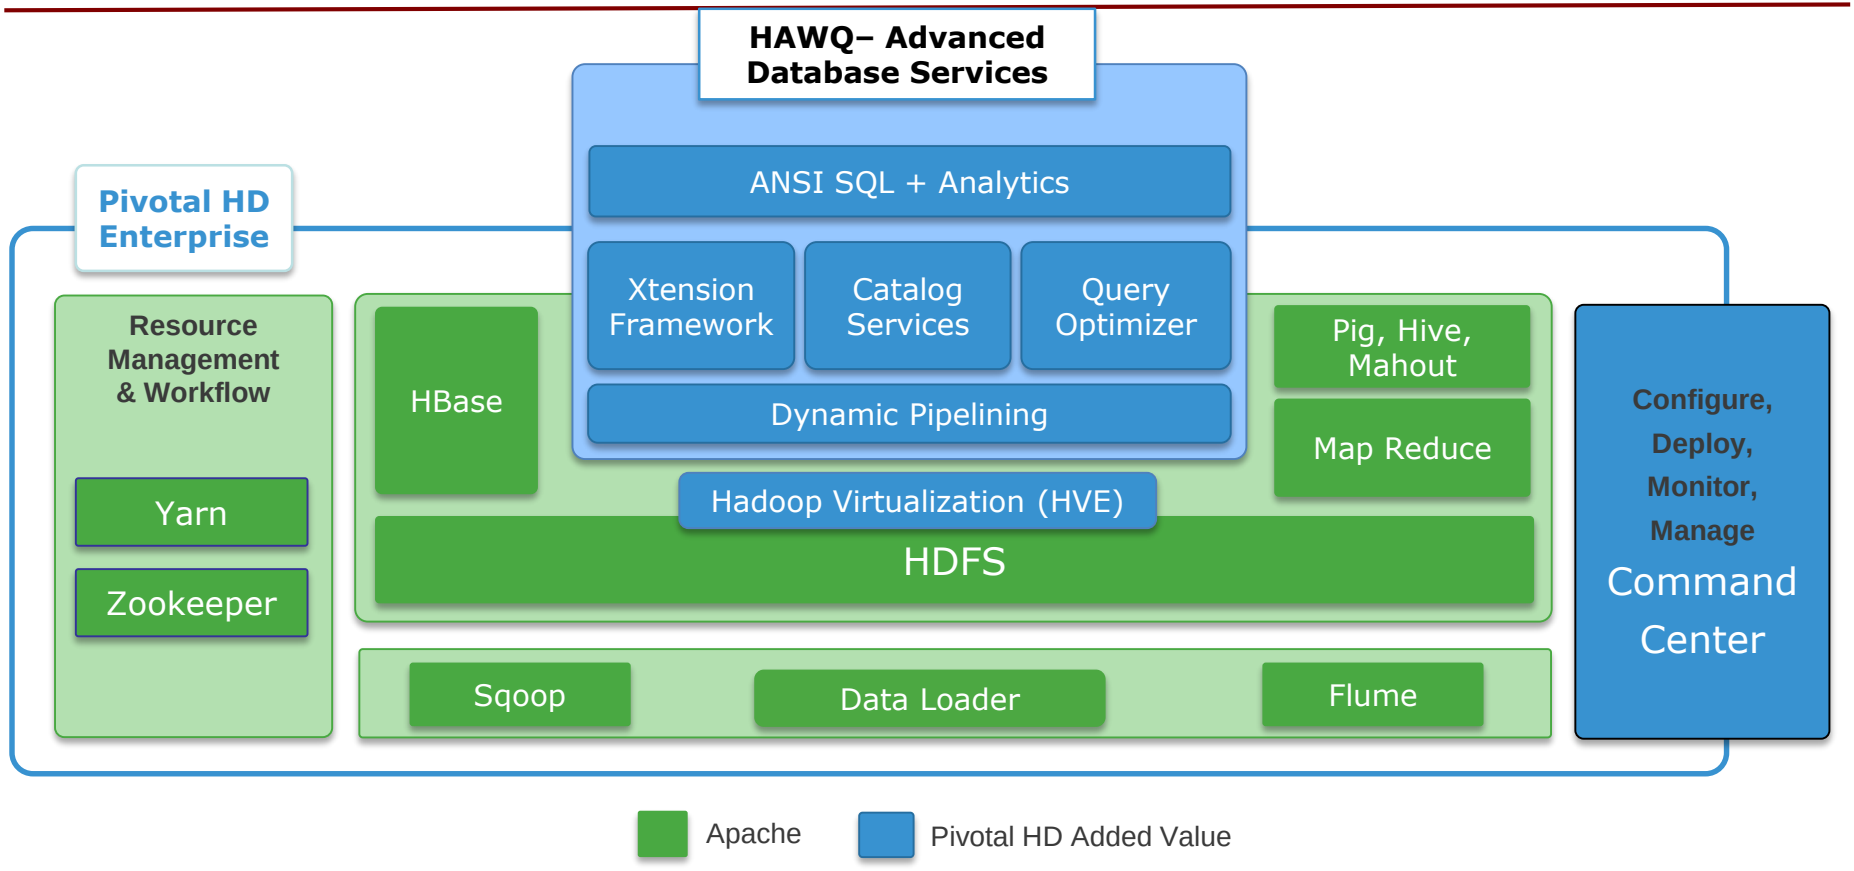
HAWQ– Advanced Database Services
ANSI SQL + Analytics
Xtension Framework
Catalog Services
Query Optimizer
Dynamic Pipelining
Pivotal HD Enterprise
Resource Management & Workflow
Yarn
Zookeeper
HBase
Hadoop Virtualization (HVE)
HDFS
Pig, Hive, Mahout
Map Reduce
Sqoop
Data Loader
Flume
Configure, Deploy, Monitor, Manage
Command Center
Apache
Pivotal HD Added Value

IBM BigInsights

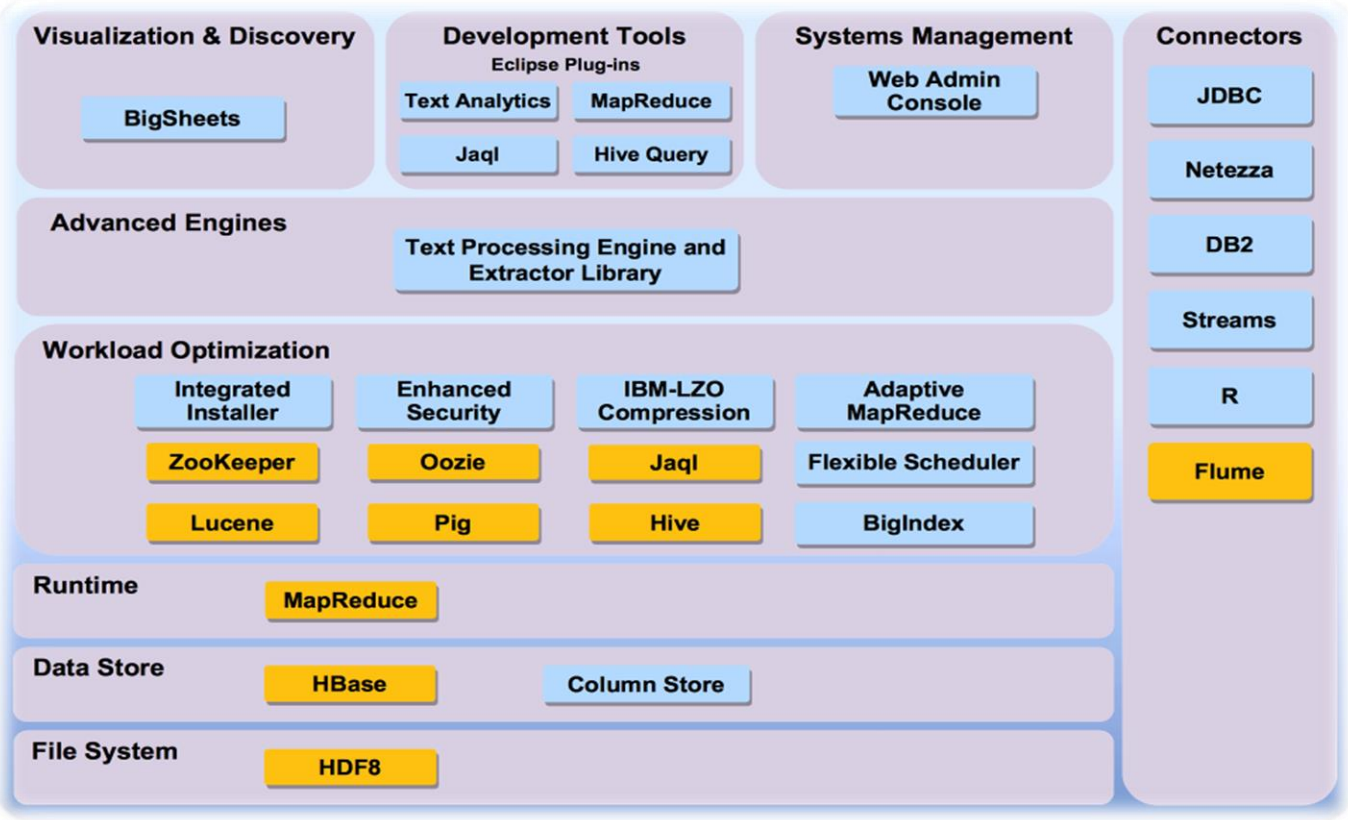

Open Data Platform

PLATINUM

THE OPEN DATA PLATFORM WILL

1. Accelerate the delivery of Big Data solutions by providing a well-defined core platform to target.
2. Define, integrate, test, and certify a standard "ODP Core" of compatible versions of select Big Data open source projects.
3. Provide a stable base against which Big Data solutions providers can qualify solutions.
4. Produce a set of tools and methods that enable members to create and test differentiated offerings based on the ODP Core.
5. Reinforce the role of the Apache Software Foundation (ASF) in the development and governance of upstream projects.
6. Contribute to ASF projects in accordance with ASF processes and Intellectual Property guidelines.
7. Support community development and outreach activities that accelerate the rollout of modern data architectures that leverage Apache Hadoop®.
8. Will help minimize the fragmentation and duplication of effort within the industry.

Ecosystem

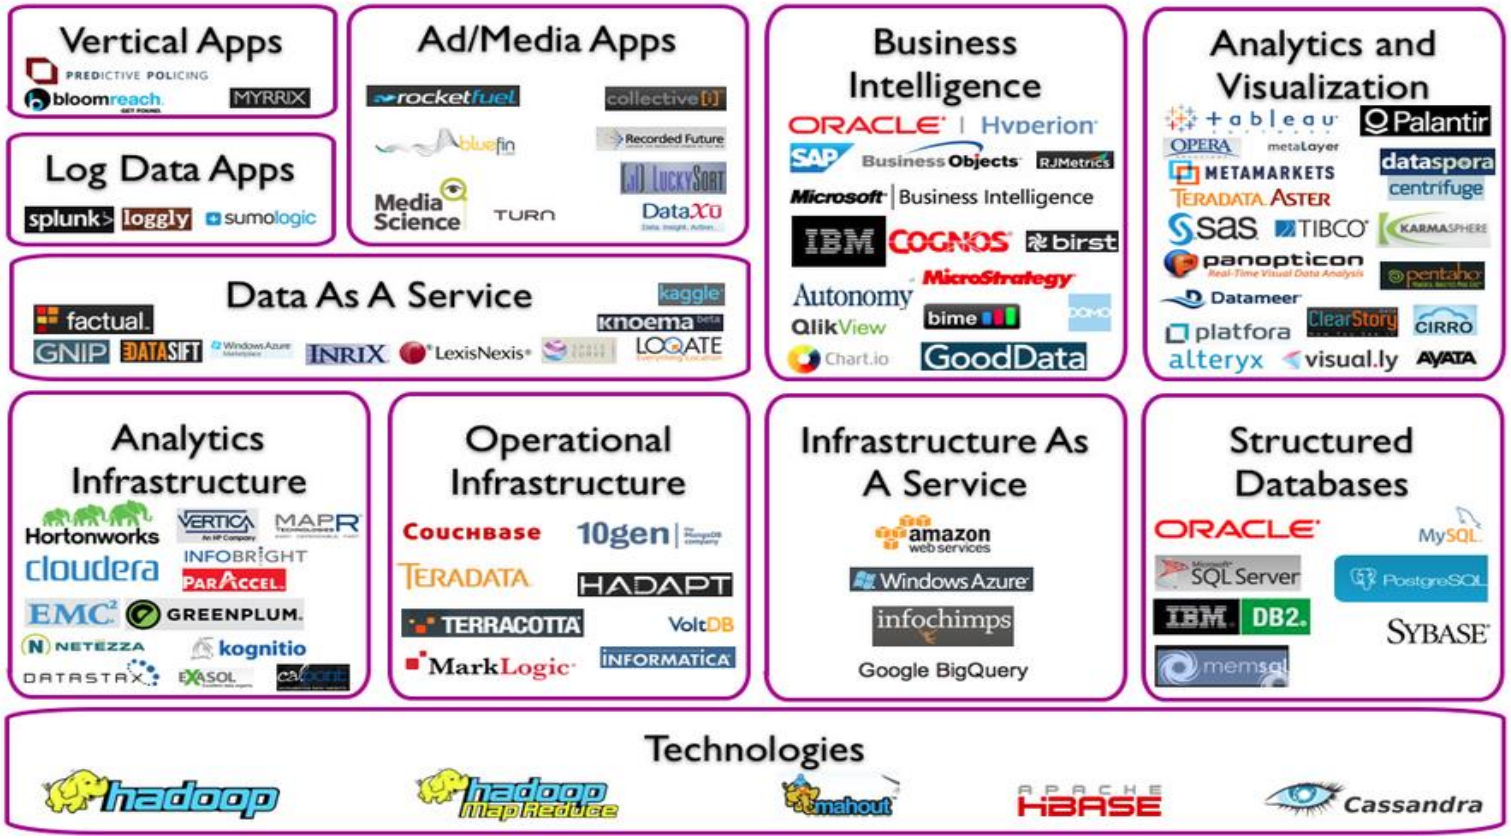

HDFS

ARCHITECTURE DEEP-DIVE

Topology of a Hadoop Cluster

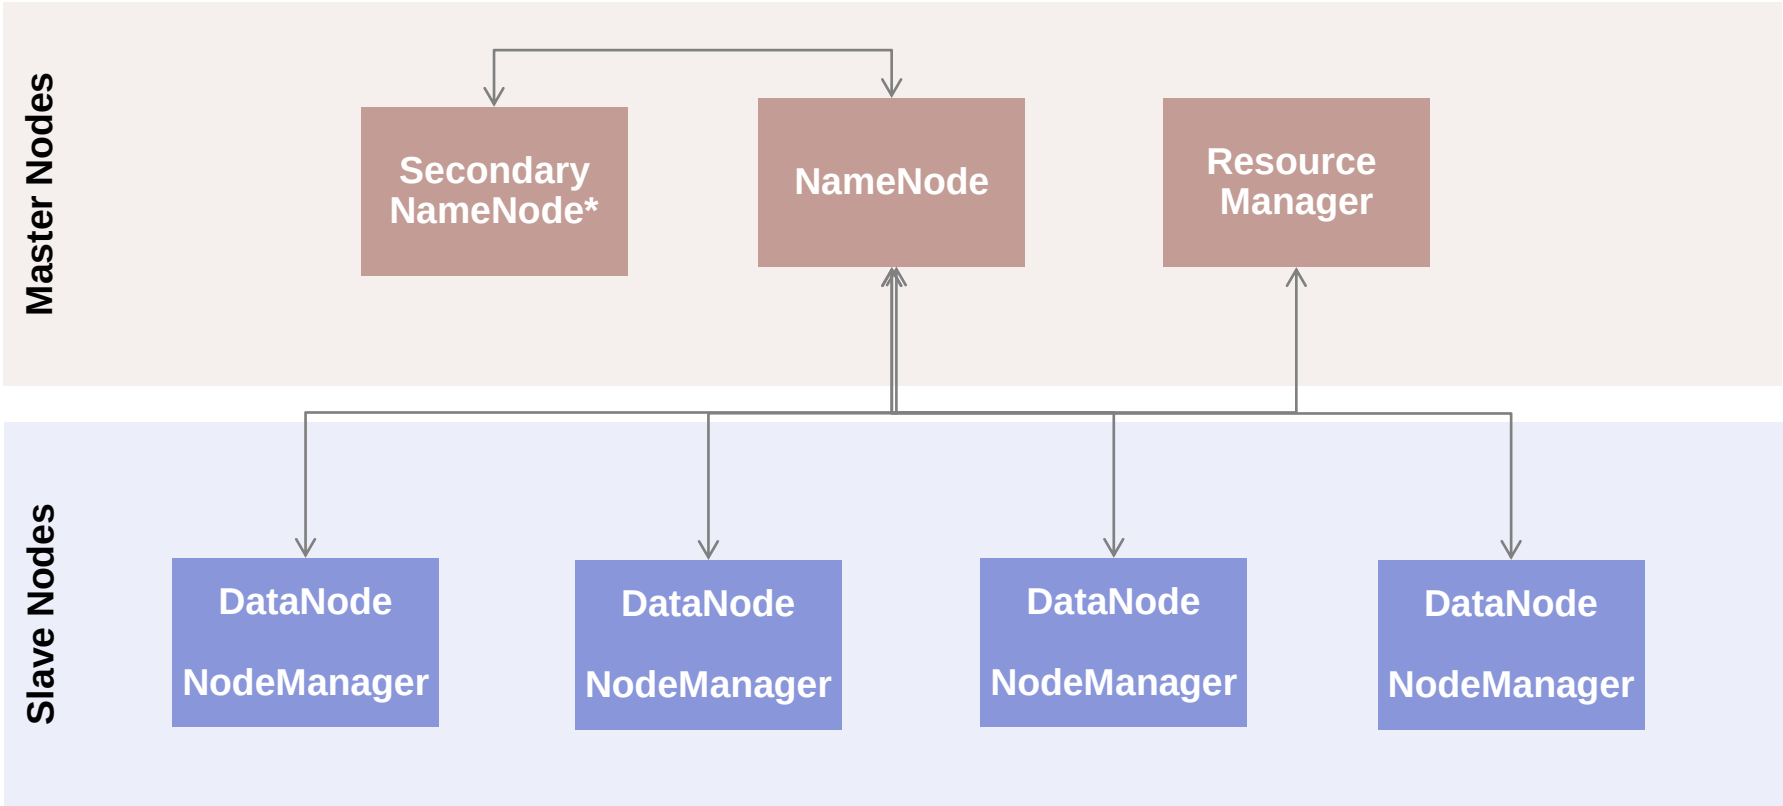

Topology of a Hadoop (MapR) Cluster

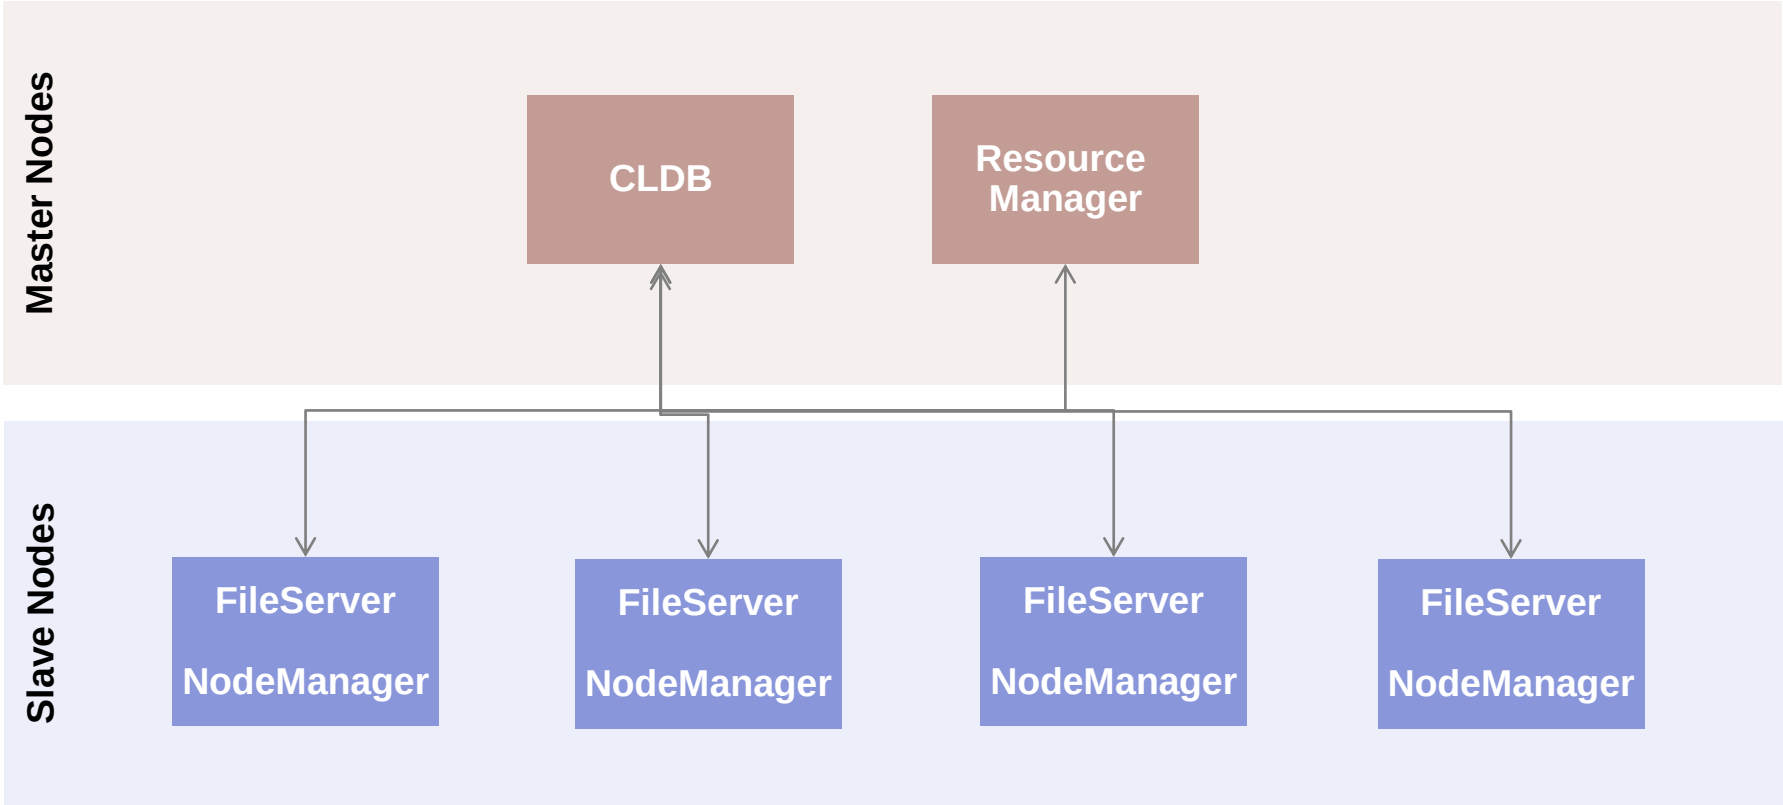

HDFS: Architecture

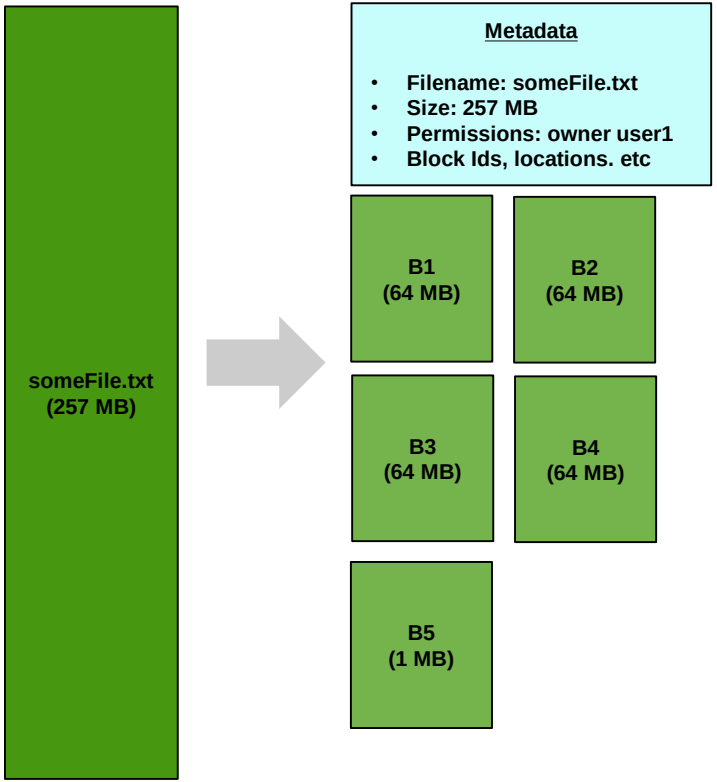

- File consists of Metadata and Data
- Data is broken into fixed sized **blocks**
- **Blocks** stored by **DataNodes**
- DataNodes
  - Host and serve blocks
  - All operations go to NN
  - No idea of files / directory / metadata
- **NameNode** holds (in memory)
  - Directory, Files Listing
  - Block replica locations
- Secondary NameNode
  - Assists NameNode

HDFS Storage Details

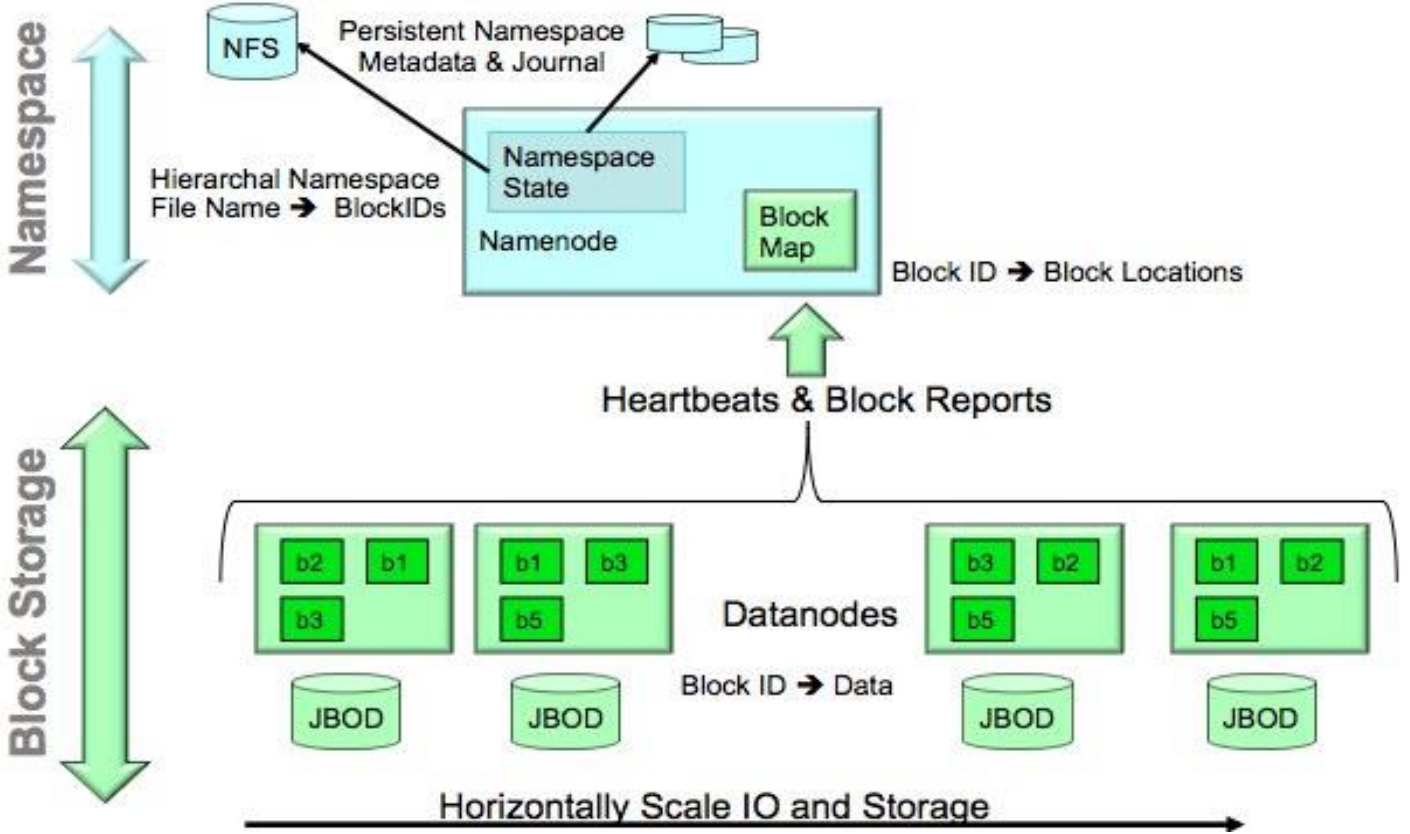

HDFS: Data Write

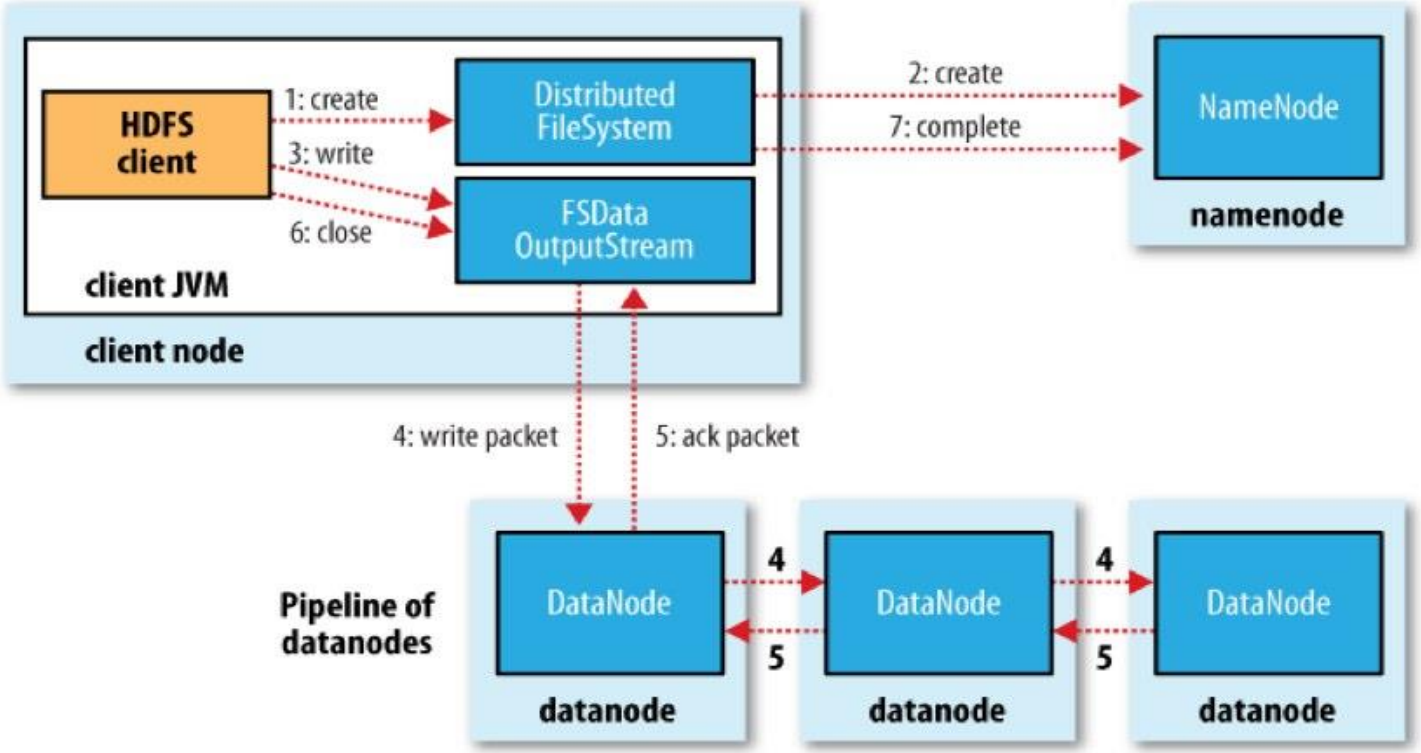

Source: Hadoop, The Definitive Guide

HDFS: Data Read

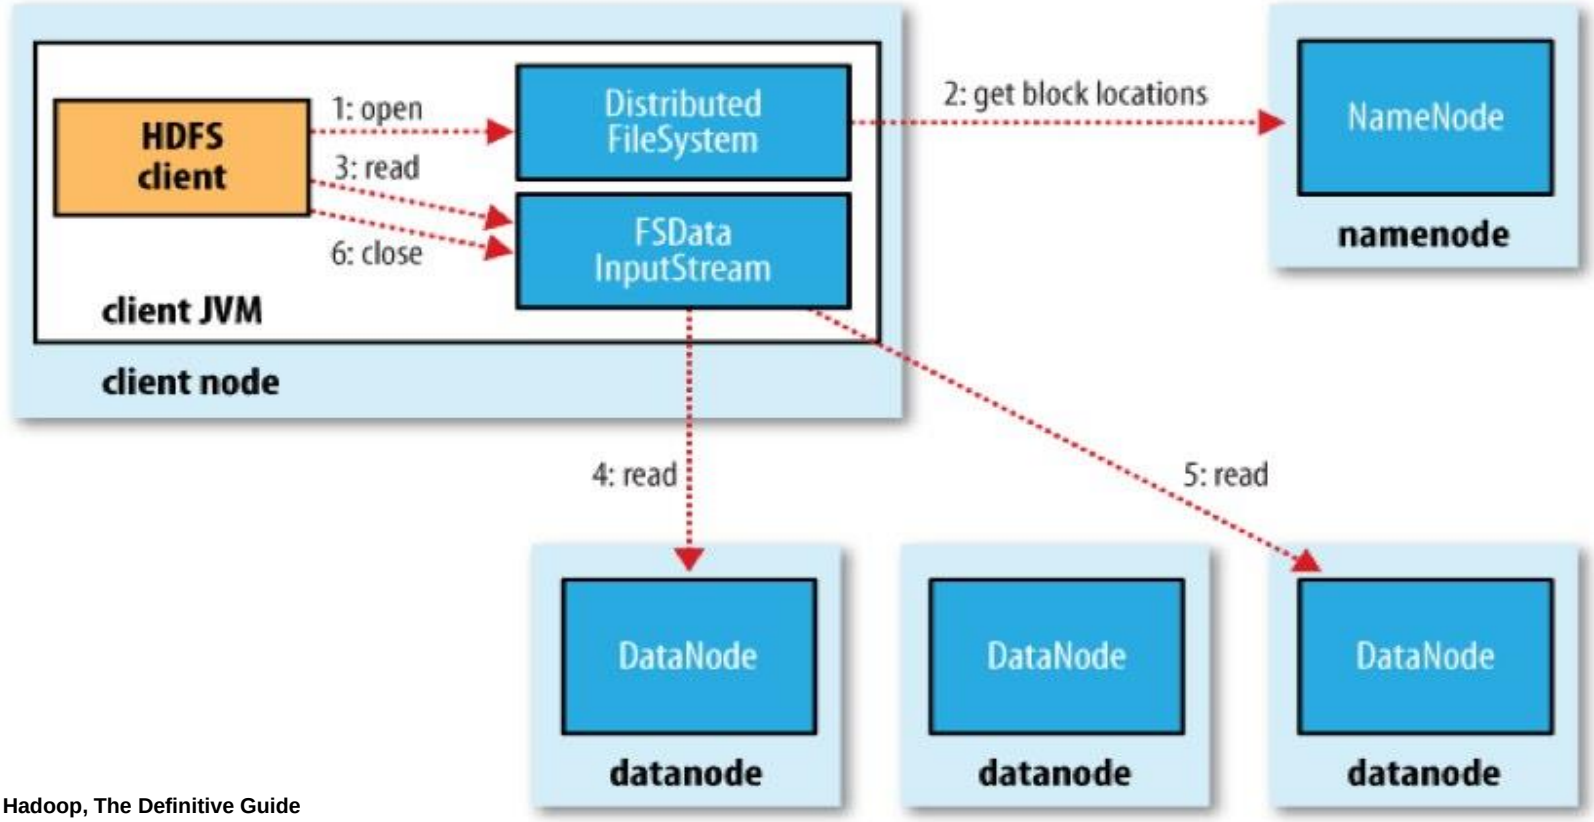

Source: Hadoop, The Definitive Guide

Alternatives to HDFS

ZALONI

- Any FileSystem that has implements the Apache Hadoop “FileSystem” API can interoperate with Hadoop
  - MapR File System (mapr-fs)
  - Amazon Simple Storage Service (S3)
  - Azure Blobstorage
  - Tachyon

MapR's Distributed NameNode

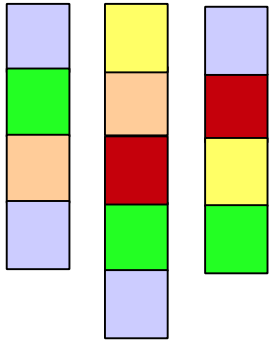

Files/directories are sharded into blocks, which are placed into containers (mini NNs) on disks

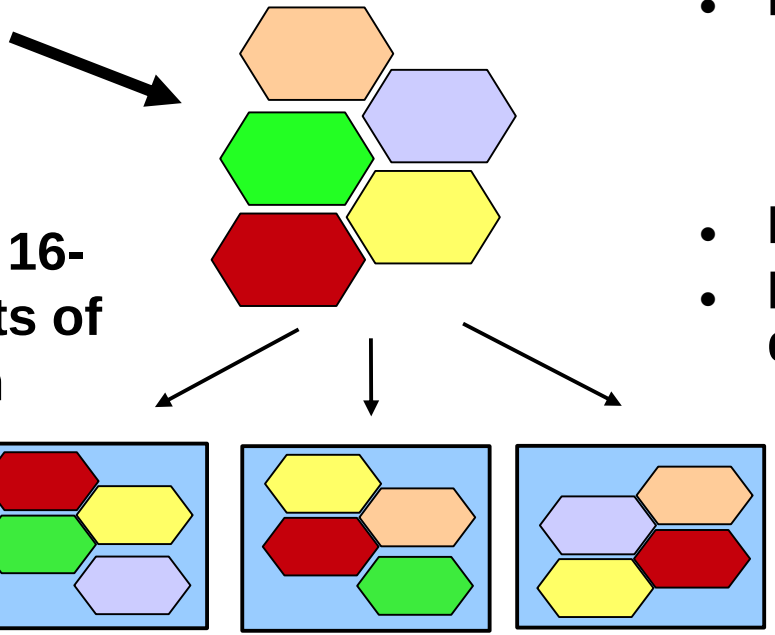

- **Each container contains**
  - **Directories & files**
  - **Data blocks**
- **Replicated on servers**
- **No need to manage directly**

Containers are 16-32 GB segments of disk, placed on nodes

http://www.slideshare.net/mcsrivas/design-scale-and-performance-of-maprs-distribution-for-hadoop

Container Location and Replication

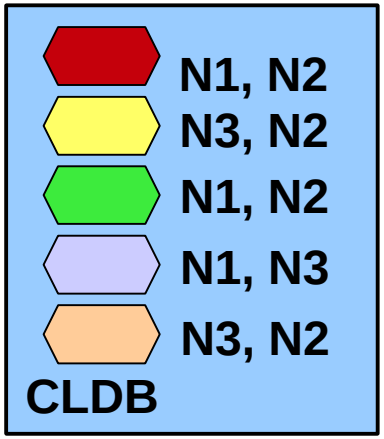

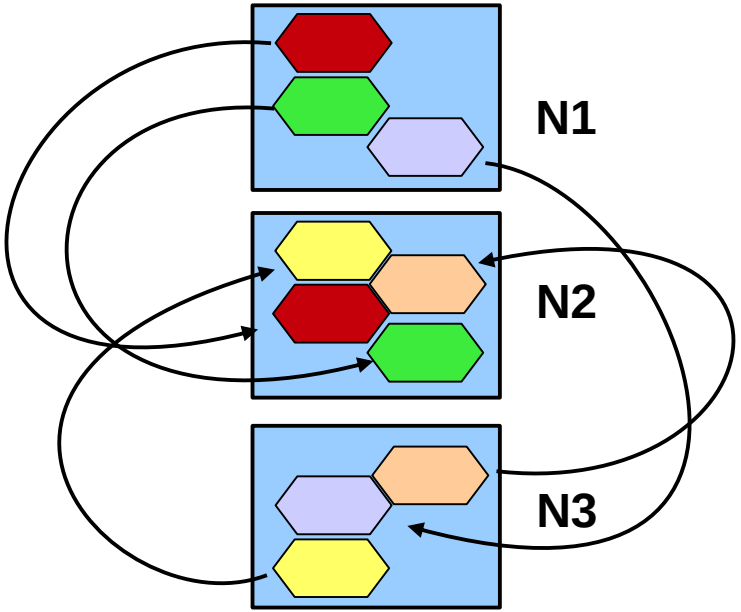

Container location database (CLDB) keeps track of nodes hosting each container and replication chain order

MapR Distributed NameNode Scaling

ZALONI

Containers represent 16 - 32GB of data

- Each can hold up to 1 Billion files and directories
- 100M containers = ~ 2 Exabytes (a very large cluster)

250 bytes DRAM to cache a container

- 25GB to cache all containers for 2EB cluster
  - But not necessary, can page to disk
- Typical large 10PB cluster needs 2GB

Container-reports are 100x - 1000x < HDFS block-reports

- Serve 100x more data-nodes
- Increase container size to 64G to serve 4EB cluster
  - Map/reduce not affected

http://www.slideshare.net/mcsrivas/design-scale-and-performance-of-maprs-distribution-for-hadoop

MapR Distributed NameNode HA vs Hadoop

ZALONI

MapR

1. apt-get install mapr-cldb
while cluster is online

Apache Hadoop*

1. Stop cluster very carefully
2. Move fs.checkpoint.dir onto NAS (eg. NetApp)
3. Install, configure DRBD + Heartbeat packages
   i. yum -y install drbd82 kmod-drbd82 heartbeat
   ii. chkconfig -add heartbeat (both machines)
   iii. edit /etc/drbd.conf on 2 machines
   iv-xxxix. make raid-0 md, ask drbd to manage raid md, zero it if drbd dies & try again
   xxxx. mkfs ext3 on it, mount /hadoop (both machines)
   xxxxi. install all rpms in /hadoop, but don't run them yet (chkconfig off)
   xxxxii. umount /hadoop (!!)
   xxxxiii. edit 3 files /etc/ha.d/* to configure heartbeat

. . .

40. Restart cluster. If any problems, start at /var/log/ha.log for hints on what went wrong.

*As described in www.cloudera.com/blog/2009/07/hadoop-ha-configuration
Author: Christophe Bisciglia, Cloudera.

http://www.slideshare.net/mcsrivas/design-scale-and-performance-of-maprs-distribution-for-hadoop

MapR Volumes

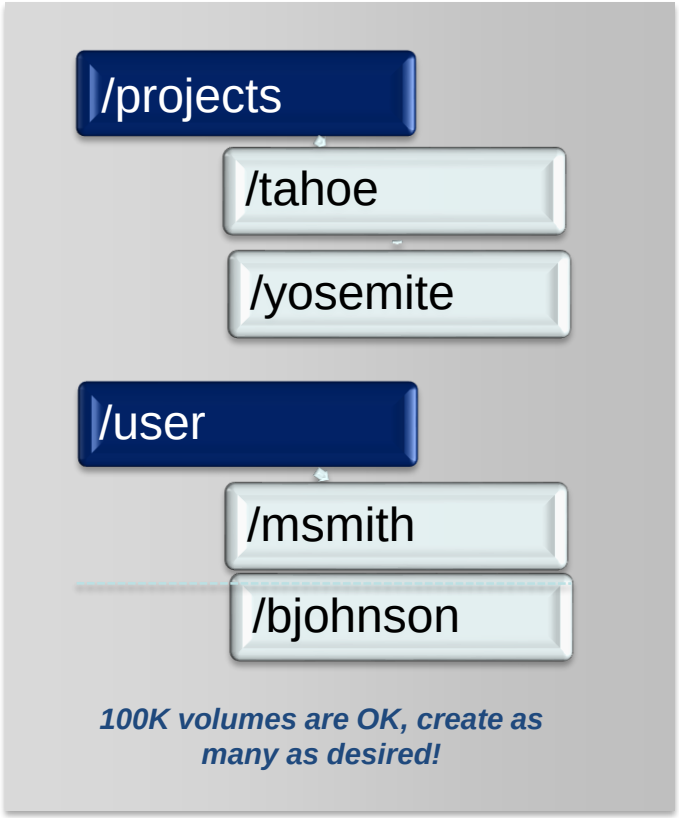

Volumes allow management attributes to be applied in a scalable way at a very granular level and with flexibility

- Replication factor
- Scheduled mirroring
- Scheduled snapshots
- Data placement control
- User access and tracking
- Administrative permissions

http://www.slideshare.net/mcsrivas/design-scale-and-performance-of-maprs-distribution-for-hadoop

MapR NFS advantage for data import/export

With MapR, use NFS

1. mount /mapr
   real-time, HA

Otherwise, use Flume/Scribe

1. Set up sinks (*find unused machines??)*
2. Set up intrusive agents
   i. tail("xxx"), tailDir("y")
   ii. agentBESink
3. All reliability levels lose data
   i. best-effort
   ii. one-shot
   iii. disk fail-over
   iv. end-to-end
4. Data not available now

http://www.slideshare.net/mcsrivas/design-scale-and-performance-of-maprs-distribution-for-hadoop

When to use HDFS

- Store large structured and unstructured files
  - Server logs
  - Relational Files
  - Data Feeds
  - Archive
  - Satellite images

When to not use HDFS

ZALONI

- HDFS is not suited for:
  - Large number of small files
    - ✓ Small files can always be concatenated
  - Files that are modified often
    - ✓ Create new files instead
  - Files that are randomly accessed
    - ✓ B-Trees, etc.

Accessing HDFS

ZALONI

- Command line tools
- Java API
- WebDFS API
- Third-party access
  - FUSE
  - Web UI
  - HDFS over FTP

HDFS Characteristics: Review

ZALONI

- Designed for **modest number of Large files** (millions instead of billions)
- **Sequential access** not Random access
- Write Once, Read Many
- Data **is split into chunks** and stored in multiple nodes as blocks
- **Namenode** maintains the block locations
- Blocks get **replicated** over the data nodes
- HDFS 2.x Features:
  - **High Availability** with Active and Standby NameNode
  - **Namespace Federation** for scalability
  - **Snapshots** to enable point-in-time recovery
  - **NFS** Gateway

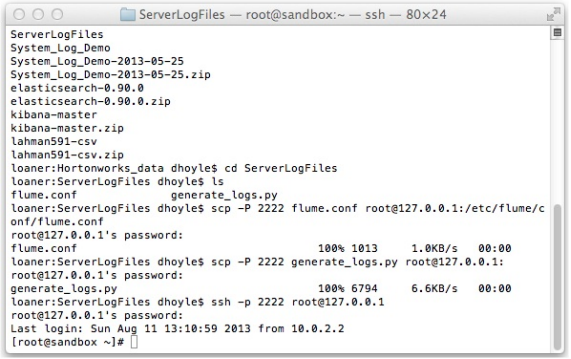

HDFS

DEMO

END OF PART 1